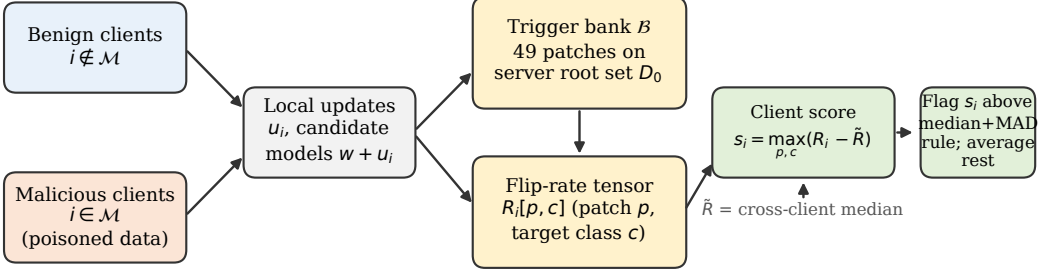
Benign clients
$i \notin \mathcal{M}$
Malicious clients
$i \in \mathcal{M}$
(poisoned data)
Local updates
u_i, candidate
models $w + u_i$
Trigger bank $\mathcal{B}$
49 patches on
server root set D_0
Flip-rate tensor
$R_i[p, c]$ (patch p,
target class c)
Client score
$s_i = \max_{p,c}(R_i - \tilde{R})$
$\tilde{R}$ = cross-client median
Flag s_i above
median+MAD
rule; average
rest